

```
In [1]: import pandas as pd
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
pd.options.display.max_rows = 999
pd.options.display.max_columns = 90

# import scikit-Learn Libraries
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn import preprocessing
```

Load the data (sql db)

```
In [2]: df_raw = pd.read_csv('../data/01_raw/bmarket.csv')
```

1. Data Understanding and Insights

```
In [3]: # Insights from running .info()
# 'Housing Loan' and 'Personal Loan' has null values

df_raw.info() # df.isna().sum()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Client ID        41188 non-null   int64  
 1   Age              41188 non-null   object  
 2   Occupation       41188 non-null   object  
 3   Marital Status   41188 non-null   object  
 4   Education Level  41188 non-null   object  
 5   Credit Default   41188 non-null   object  
 6   Housing Loan     16399 non-null   object  
 7   Personal Loan    37042 non-null   object  
 8   Contact Method   41188 non-null   object  
 9   Campaign Calls   41188 non-null   int64  
 10  Previous Contact Days 41188 non-null   int64  
 11  Subscription Status 41188 non-null   object  
dtypes: int64(3), object(9)
memory usage: 3.8+ MB
```

```
In [4]: # Check if there are any duplicated client ID

duplicated = df_raw['Client ID'].duplicated().any()
print(duplicated)
```

```
False
```

```
In [5]: # Observation from 'Age'
df_raw['Age'].value_counts()
```

```
Out[5]: Age
150 years    4197
31 years     1747
32 years     1646
33 years     1643
36 years     1606
35 years     1584
34 years     1577
30 years     1536
37 years     1345
29 years     1310
38 years     1283
39 years     1281
41 years     1142
40 years     1027
42 years     1020
45 years      979
43 years      950
46 years      920
44 years      892
28 years      892
48 years      875
47 years      833
50 years      778
27 years      765
49 years      758
52 years      689
51 years      676
53 years      654
56 years      636
26 years      630
54 years      628
57 years      597
55 years      589
25 years      536
58 years      511
24 years      414
59 years      409
60 years      254
23 years      207
22 years      123
21 years       90
61 years       66
20 years       62
62 years       55
64 years       55
66 years       54
63 years       50
71 years       47
65 years       42
70 years       40
19 years       38
73 years       33
69 years       32
68 years       31
76 years       30
74 years       30
72 years       30
80 years       29
18 years       27
78 years       23
67 years       21
75 years       20
88 years       19
77 years       18
81 years       16
83 years       16
82 years       16
85 years       15
79 years       13
86 years        8
84 years        7
17 years        5
92 years        4
89 years        2
91 years        2
98 years        2
95 years        1
Name: count, dtype: int64
```

```
In [6]: # Observation from 'Occupation'  
df_raw['Occupation'].value_counts()
```

```
Out[6]: Occupation  
admin.          10422  
blue-collar     9254  
technician      6743  
services        3969  
management      2924  
retired         1720  
entrepreneur    1456  
self-employed    1421  
housemaid       1060  
unemployed      1014  
student          875  
unknown          330  
Name: count, dtype: int64
```

```
In [7]: # Observations from 'Marital Status'  
df_raw['Marital Status'].value_counts()
```

```
Out[7]: Marital Status  
married         24928  
single          11568  
divorced        4612  
unknown          80  
Name: count, dtype: int64
```

```
In [8]: # Observation from 'Education Level'  
df_raw['Education Level'].value_counts()
```

```
Out[8]: Education Level  
university.degree   12168  
high.school         9515  
basic.9y            6045  
professional.course 5243  
basic.4y            4176  
basic.6y            2292  
unknown             1731  
illiterate           18  
Name: count, dtype: int64
```

```
In [9]: # Observation from 'Credit Default'  
df_raw['Credit Default'].value_counts()
```

```
Out[9]: Credit Default  
no            32588  
unknown        8597  
yes            3  
Name: count, dtype: int64
```

```
In [10]: # Observation from 'Housing Loan'  
df_raw['Housing Loan'].value_counts()
```

```
Out[10]: Housing Loan  
yes            8595  
no             7411  
unknown        393  
Name: count, dtype: int64
```

```
In [11]: # Observation from 'Personal Loan'  
df_raw['Personal Loan'].value_counts()
```

```
Out[11]: Personal Loan  
no            30532  
yes           5633  
unknown        877  
Name: count, dtype: int64
```

```
In [12]: # Observation from 'Contact Method'  
# Repeat in contact method  
df_raw['Contact Method'].value_counts()
```

```
Out[12]: Contact Method  
Cell           13100  
cellular       13044  
Telephone      7585  
telephone      7459  
Name: count, dtype: int64
```

```
In [13]: # Observation from 'Campaign Calls'  
# There are negative values  
df_raw['Campaign Calls'].value_counts()
```

```
Out[13]: Campaign Calls
```

1	15874
2	9446
3	4807
4	2405
-1	1768
5	1451
-2	1124
6	893
7	566
-3	534
8	365
9	253
-4	246
10	206
11	156
-5	148
12	107
-6	86
13	84
-7	63
14	58
17	51
15	47
16	46
-8	35
-9	30
18	30
20	27
19	24
21	22
-11	21
-10	19
-12	18
24	15
22	14
27	11
-14	11
23	11
29	9
-13	8
26	8
30	7
31	7
-17	7
25	7
28	6
-23	5
-16	5
-15	4
33	4
34	3
32	3
-18	3
-22	3
-20	3
35	3
42	2
-28	2
43	2
40	2
-21	2
-35	2
-19	2
-29	1
-32	1
56	1
-41	1
39	1
37	1
-25	1

```
Name: count, dtype: int64
```

```
In [14]: # Observation from 'Previous Contact Days'  
df_raw['Previous Contact Days'].value_counts()
```

```
Out[14]: Previous Contact Days
999    39673
3      439
6      412
4     118
9      64
2      61
7      60
12     58
10     52
5      46
13     36
11     28
1      26
15     24
14     20
8      18
0      15
16     11
17      8
18      7
19      3
22      3
21      2
26      1
27      1
20      1
25      1
Name: count, dtype: int64
```

```
In [15]: # Observation from 'Subscription Status'
# The number of non-subscription is almost 8 times more than subscriptions
df_raw['Subscription Status'].value_counts()
```

```
Out[15]: Subscription Status
no    36548
yes   4640
Name: count, dtype: int64
```

2. Data Cleaning

Since there are no duplicated Client IDs, we will not be removing any rows from the DB.

After that, we will confirm that the DB has no remaining null values by using '.isna().any()' '`.isna`' is used to identify missing data in a database.
`.any()` would return one value for each column. *TRUE* if any value in that column is *TRUE*, and *FALSE* if otherwise.
Then, we do a check to ensure the number of values match for all the columns.

```
In [16]: df_cleaned = df_raw.copy()

df_cleaned = df_raw.fillna({
    'Housing Loan': 'unknown',
    'Personal Loan': 'unknown'
})

display(df_cleaned.isna().any())
print('\n')
display(df_cleaned.info())
```

Client ID	False
Age	False
Occupation	False
Marital Status	False
Education Level	False
Credit Default	False
Housing Loan	False
Personal Loan	False
Contact Method	False
Campaign Calls	False
Previous Contact Days	False
Subscription Status	False
dtype: bool	

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Client ID        41188 non-null   int64  
 1   Age              41188 non-null   object  
 2   Occupation       41188 non-null   object  
 3   Marital Status   41188 non-null   object  
 4   Education Level  41188 non-null   object  
 5   Credit Default   41188 non-null   object  
 6   Housing Loan     41188 non-null   object  
 7   Personal Loan    41188 non-null   object  
 8   Contact Method   41188 non-null   object  
 9   Campaign Calls   41188 non-null   int64  
 10  Previous Contact Days 41188 non-null   int64  
 11  Subscription Status 41188 non-null   object  
dtypes: int64(3), object(9)
memory usage: 3.8+ MB
None

```

In 'Contact Method', some values appear different but have the same meaning.

1. 'Cell' / 'cellular'
2. 'Telephone' / 'telephone'

Therefore, we will change first value to match the second value using '.replace()'

```
df['<Column_Name>'] = df['<Column_Name>'].replace('First_value', 'Second value')
```

With this, any cells with the first value 'Cell' will be replaced with 'cellular'.

```
In [17]: df_cleaned['Contact Method'] = df_cleaned['Contact Method'].replace('Cell', 'cellular')
df_cleaned['Contact Method'] = df_cleaned['Contact Method'].replace('Telephone', 'telephone')

df_cleaned['Contact Method'].value_counts()
```

```
Out[17]: Contact Method
cellular      26144
telephone     15044
Name: count, dtype: int64
```

There is a large percentage of users that are '150 years' old.

Therefore, we have decided to remove the data entries as they are unrealistic for humans and would mislead the analysis and confuse the model.

This helps the data be more reliable and meaningful.

```
In [18]: df_cleaned['Age'] = (
    df_cleaned['Age']
    .astype(str)
    .str.replace(' years', '')
    .astype(int)
)
df_cleaned = df_cleaned.drop(df_cleaned[df_cleaned['Age'] == 150].index)

df_cleaned['Age'].value_counts()
```

```
Out[18]: Age
31    1747
32    1646
33    1643
36    1686
35    1584
34    1577
30    1536
37    1345
29    1310
38    1283
39    1281
41    1142
40    1027
42    1020
45    979
43    950
46    920
44    892
28    892
48    875
47    833
50    778
27    765
49    758
52    689
51    676
53    654
56    636
26    630
54    628
57    597
55    589
25    536
58    511
24    414
59    409
60    254
23    207
22    123
21    90
61    66
20    62
62    55
64    55
66    54
63    50
71    47
65    42
70    40
19    38
73    33
69    32
68    31
76    30
74    30
72    30
80    29
18    27
78    23
67    21
75    20
88    19
77    18
81    16
83    16
82    16
85    15
79    13
86    8
84    7
17    5
92    4
89    2
91    2
98    2
95    1
Name: count, dtype: int64
```

Some entries in 'Campaign Calls' appeared as negative numbers. However, the number of calls made to a client cannot be negative in real life. This mistake is almost certainly caused by a data entry mistake or a formatting issue.

As the actual value should logically be the positive number, we corrected these entries by removing the negative sign.

```
In [19]: df_cleaned['Campaign Calls'] = df_cleaned['Campaign Calls'].astype(str)
df_cleaned['Campaign Calls'] = df_cleaned['Campaign Calls'].str.replace('-', '')
df_cleaned['Campaign Calls'] = df_cleaned['Campaign Calls'].astype(int)

df_cleaned['Campaign Calls'].value_counts()
```

```
Out[19]: Campaign Calls
1      15799
2      9532
3      4779
4      2404
5      1426
6      878
7      567
8      361
9      255
10     201
11     159
12     116
13     84
14     61
17     52
16     48
15     43
18     31
20     29
19     23
21     22
23     16
22     15
24     13
29     10
27     10
25     8
26     8
28     7
30     6
31     6
33     4
35     4
32     3
43     2
40     2
34     2
42     1
56     1
41     1
39     1
37     1
Name: count, dtype: int64
```

```
In [20]: df_cleaned['WasContactedBefore'] = (df_cleaned['Previous Contact Days'] != 999).astype(int)
df_cleaned['PreviouslyContacted'] = df_cleaned['Previous Contact Days'].replace({999: 0})
```

```
In [21]: bins = [1, 25, 35, 45, 55, 65, 120]
labels = ['1-25', '26-35', '36-45', '46-55', '56-65', '66+']

# create age_group in df_filled
df_cleaned['Age Group'] = pd.cut(df_cleaned['Age'],
                                   bins=bins,
                                   labels=labels,
                                   right=False)
```

```
In [22]: # data cleaning done, display overview
df_cleaned.to_csv("df_cleaned.csv", index=False)
df_cleaned.head()
```

Out[22]:

	Client ID	Age	Occupation	Marital Status	Education Level	Credit Default	Housing Loan	Personal Loan	Contact Method	Campaign Calls	Previous Contact Days	Subscription Status
0	32885	57	technician	married	high.school	no	no	yes	cellular	1	999	n
1	3170	55	unknown	married	unknown	unknown	yes	no	telephone	2	999	n
2	32207	33	blue-collar	married	basic.9y	no	no	no	cellular	1	999	n
3	9404	36	admin.	married	high.school	no	no	no	telephone	4	999	n
4	14021	27	housemaid	married	high.school	no	unknown	no	cellular	2	999	n

Data Analysis

Financial Loans Analysis

```
In [23]: # Find out the number of people with loans or with credit default
cd = df_cleaned['Credit Default'].value_counts()
pl = df_cleaned['Personal Loan'].value_counts()
hl = df_cleaned['Housing Loan'].value_counts()
sub = df_cleaned['Subscription Status'].value_counts()
print(cd)
print(pl)
print(hl)
print(sub)
```

Credit Default
no 29272
unknown 7716
yes 3
Name: count, dtype: int64
Personal Loan
no 27427
yes 5041
unknown 4523
Name: count, dtype: int64
Housing Loan
unknown 22659
yes 7688
no 6644
Name: count, dtype: int64
Subscription Status
no 32800
yes 4191
Name: count, dtype: int64

Areas of Analysis:

1. Percentage of people without credit default, with housing loan or personal loan and are subscribed
2. Combination of housing and personal loan with the subscription status

1. Analysis of housing loan with subscription status

```
In [24]: #Defining all financial loans
fill_cols = ['Housing Loan', 'Personal Loan', 'Credit Default']
target_col = 'Subscription Status'
feature_col_hl = 'Housing Loan'
feature_col_pl = 'Personal Loan'
feature_col_cd = 'Credit Default'

#Groups housing loan data by the status
housing_loan_analysis = df_cleaned.groupby([feature_col_hl, target_col]).size().unstack(fill_value=0)
housing_loan_analysis['Total'] = housing_loan_analysis.sum(axis=1)
#Calculation of percentage subscribed
if 'yes' in housing_loan_analysis.columns:
    housing_loan_analysis['Subscription Rate (%)'] = (housing_loan_analysis['yes'] / housing_loan_analysis['Total']) * 100
else:
    housing_loan_analysis['Subscription Rate (%)'] = 0
housing_loan_analysis = housing_loan_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(housing_loan_analysis)

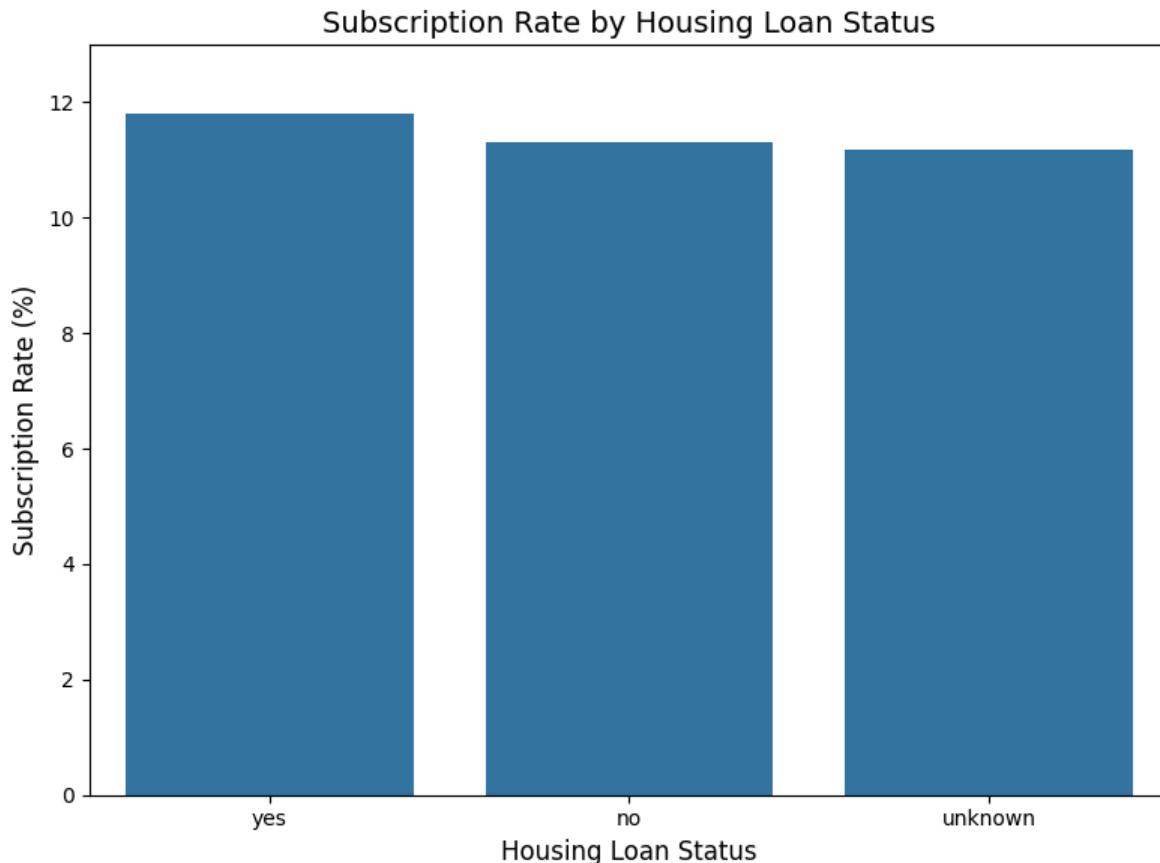
# Plotting bar graph
```

```

plt.figure(figsize=(8, 6))
sns.barplot(x=housing_loan_analysis.index, y='Subscription Rate (%)', data=housing_loan_analysis)
plt.title(f'Subscription Rate by {feature_col_hl} Status', fontsize=14)
plt.xlabel(f'{feature_col_hl} Status', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, housing_loan_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('housing_loan_subscription_rate.png')
plt.show()

```

Subscription Status	no	yes	Total	Subscription Rate (%)
Housing Loan				
yes	6780	908	7688	11.810614
no	5893	751	6644	11.303432
unknown	20127	2532	22659	11.174368



2. Analysis of personal loan with subscription status

```

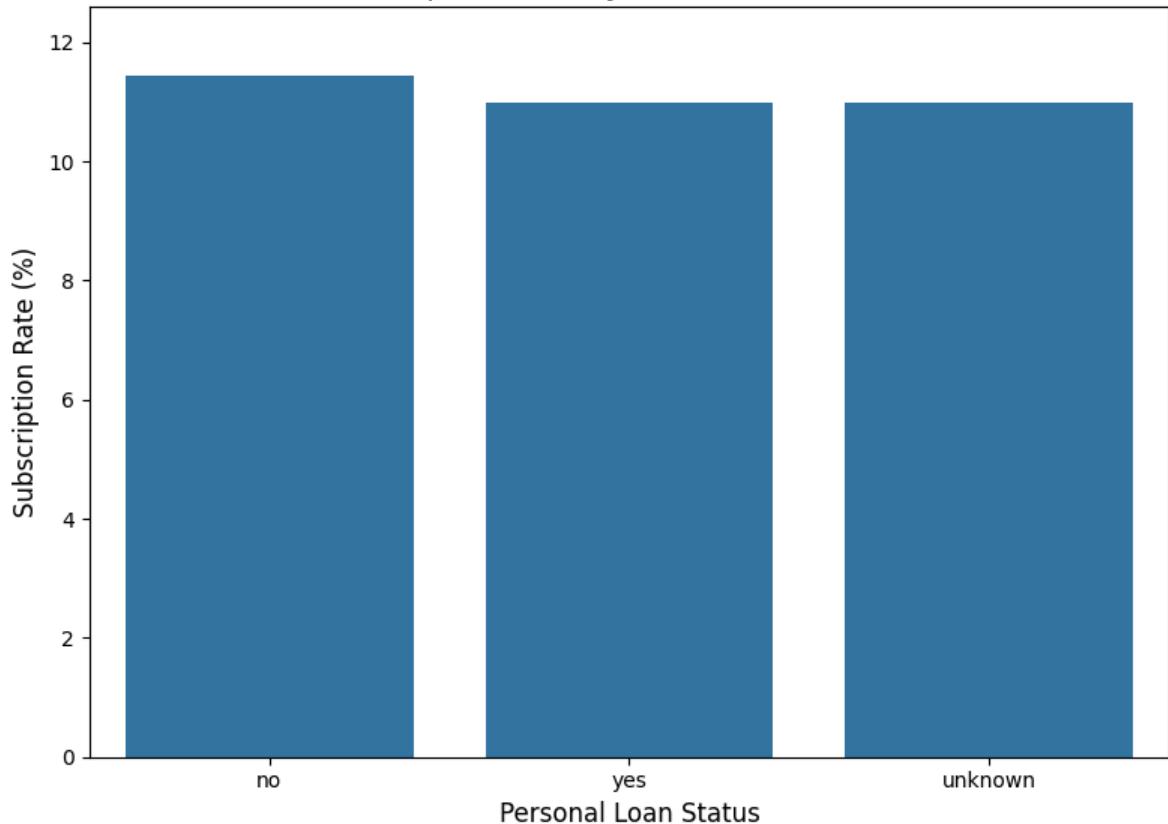
In [25]: personal_loan_analysis = df_cleaned.groupby([feature_col_pl, target_col]).size().unstack(fill_value=0)
personal_loan_analysis['Total'] = personal_loan_analysis.sum(axis=1)
if 'yes' in personal_loan_analysis.columns:
    personal_loan_analysis['Subscription Rate (%)'] = (personal_loan_analysis['yes'] / personal_loan_analysis['Total'])
else:
    personal_loan_analysis['Subscription Rate (%)'] = 0
personal_loan_analysis = personal_loan_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(personal_loan_analysis)

# Plotting the bar graph
plt.figure(figsize=(8, 6))
sns.barplot(x=personal_loan_analysis.index, y='Subscription Rate (%)', data=personal_loan_analysis)
plt.title(f'Subscription Rate by {feature_col_pl} Status', fontsize=14)
plt.xlabel(f'{feature_col_pl} Status', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, personal_loan_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('personal_loan_subscription_rate.png')
plt.show()

```

Subscription Status	no	yes	Total	Subscription Rate (%)
Personal Loan				
no	24287	3140	27427	11.448573
yes	4487	554	5041	10.989883
unknown	4026	497	4523	10.988282

Subscription Rate by Personal Loan Status

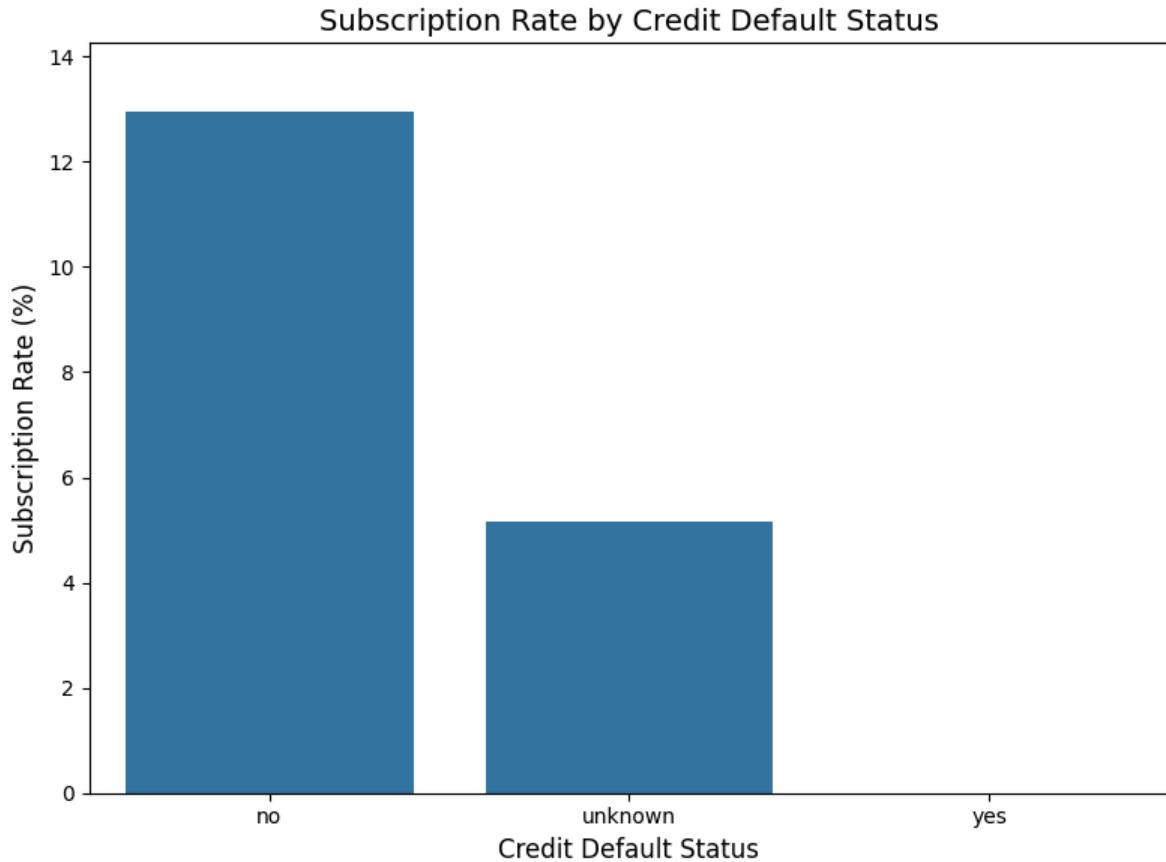


3. Analysis of credit default with subscription status

```
In [26]: credit_default_analysis = df_cleaned.groupby([feature_col_cd, target_col]).size().unstack(fill_value=0)
credit_default_analysis['Total'] = credit_default_analysis.sum(axis=1)
if 'yes' in credit_default_analysis.columns:
    credit_default_analysis['Subscription Rate (%)'] = (credit_default_analysis['yes'] / credit_default_analysis['Total'])
else:
    credit_default_analysis['Subscription Rate (%)'] = 0
credit_default_analysis = credit_default_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(credit_default_analysis)

# Plotting the bar graph
plt.figure(figsize=(8, 6))
sns.barplot(x=credit_default_analysis.index, y='Subscription Rate (%)', data=credit_default_analysis)
plt.title(f'Subscription Rate by {feature_col_cd} Status', fontsize=14)
plt.xlabel(f'{feature_col_cd} Status', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, credit_default_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('credit_default_subscription_rate.png')
plt.show()
```

Subscription Status	no	yes	Total	Subscription Rate (%)
Credit Default				
no	25479	3793	29272	12.957775
unknown	7318	398	7716	5.158113
yes	3	0	3	0.000000



4. Combining the loans to find out the relationship between each combination of housing and personal loan with the subscription status

```
In [27]: df_cleaned['Combined Loan Status'] = df_cleaned['Housing Loan'] + ' & ' + df_cleaned['Personal Loan']
feature_col_cl = 'Combined Loan Status'
print(f"4. {feature_col_cl} (Housing & Personal) vs {target_col} Analysis")

combined_loan_analysis = df_cleaned.groupby([feature_col_cl, target_col]).size().unstack(fill_value=0)
combined_loan_analysis['Total'] = combined_loan_analysis.sum(axis=1)
if 'yes' in combined_loan_analysis.columns:
    combined_loan_analysis['Subscription Rate (%)'] = (combined_loan_analysis['yes'] / combined_loan_analysis['Total'])
else:
    combined_loan_analysis['Subscription Rate (%)'] = 0
combined_loan_analysis = combined_loan_analysis.sort_values(by='Subscription Rate (%)')
print(combined_loan_analysis)

# Plotting
plt.figure(figsize=(12, 6))
sns.barplot(x=combined_loan_analysis.index, y='Subscription Rate (%)', data=combined_loan_analysis)
plt.title(f'{feature_col_cl} (Housing & Personal) vs Subscription Status Analysis', fontsize=14)
plt.xlabel(f'{feature_col_cl} (Housing & Personal)', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, combined_loan_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.savefig('combined_loan_subscription_rate.png')
plt.show()
```

4. Combined Loan Status (Housing & Personal) vs Subscription Status Analysis

Subscription Status	no	yes	Total	Subscription Rate (%)
Combined Loan Status				
no & yes	744	76	820	9.268293
unknown & unknown	2733	322	3055	10.540098
unknown & yes	2697	339	3036	11.166008
unknown & no	14697	1871	16568	11.292854
no & no	4536	590	5126	11.509949
yes & unknown	680	90	770	11.688312
yes & yes	1046	139	1185	11.729958
yes & no	5054	679	5733	11.843712
no & unknown	613	85	698	12.177650



Observations:

Even though there are many users whose credit default, housing loan and/or personal loan are unknown, they will still be used in the analysis as they contain data in other columns.

For credit default and loans:

- People with No Personal Loans are MORE LIKELY to subscribe compared to those who have
- People with No Housing Loans are LESS LIKELY to subscribe compared to those who have
- People with No Credit Default are ALMOST CERTAIN to subscribe compared to those who have

As only 3 users have credit default and none of them are subscribed compared to 10.25% of users with no credit default, credit default is an essential indicator to whether a user is more likely to subscribe.

For the combined housing and personal loan, excluding the combinations with unknowns (provides no information on the client), there is a higher percentage for those with housing loan and no personal loan followed by those with no housing and personal loan

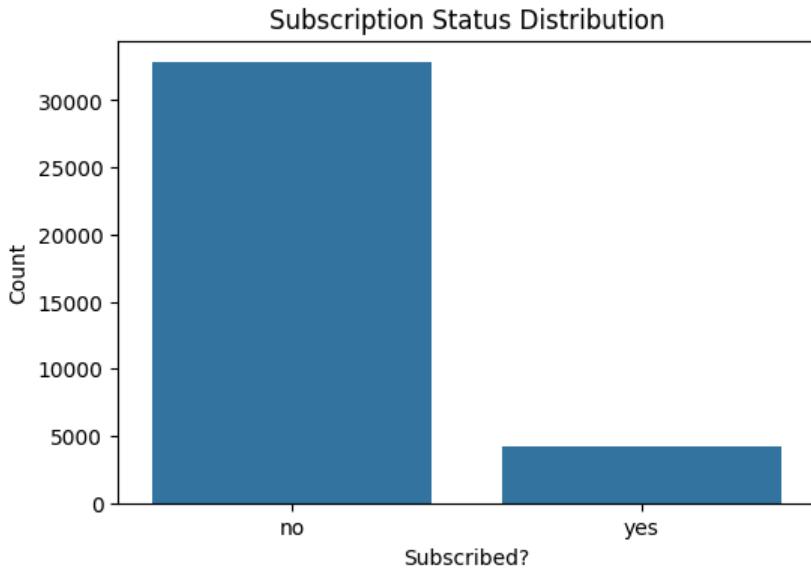
Subscription status distribution

```
In [28]: sub_counts = df_cleaned['Subscription Status'].value_counts()

plt.figure(figsize=(6,4))
sns.countplot(data=df_cleaned, x='Subscription Status')
plt.title("Subscription Status Distribution")
plt.xlabel("Subscribed?")
plt.ylabel("Count")
plt.show()

# Percentage calculation
yes_rate = (sub_counts.get('yes', 0) / df_cleaned.shape[0]) * 100
no_rate = (sub_counts.get('no', 0) / df_cleaned.shape[0]) * 100

yes_rate, no_rate
```



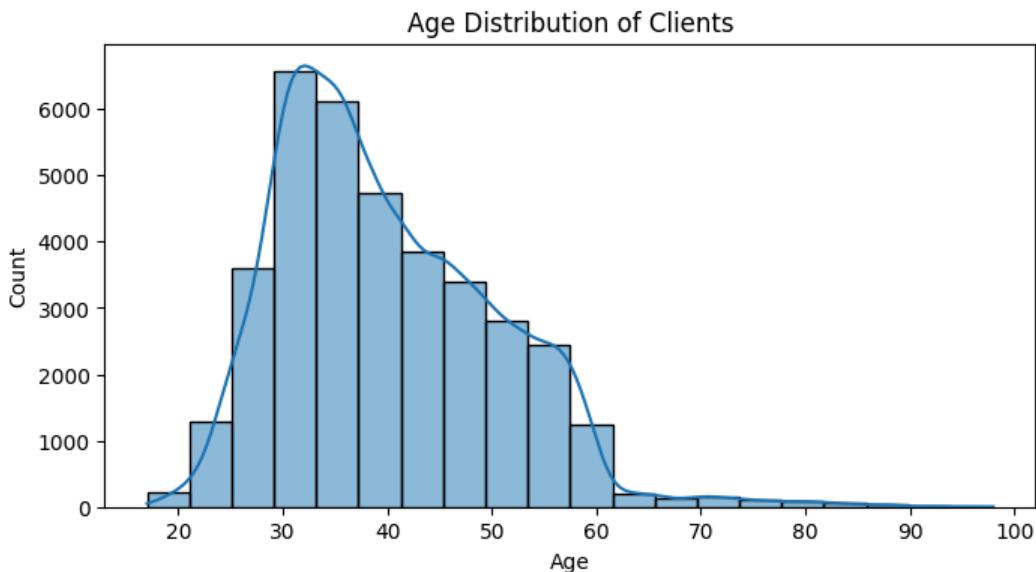
```
Out[28]: (np.float64(11.329782920169771), np.float64(88.67021707983022))
```

Most clients did not subscribe to the term deposit.

Only a small portion of clients said "yes", making this a highly imbalanced dataset.

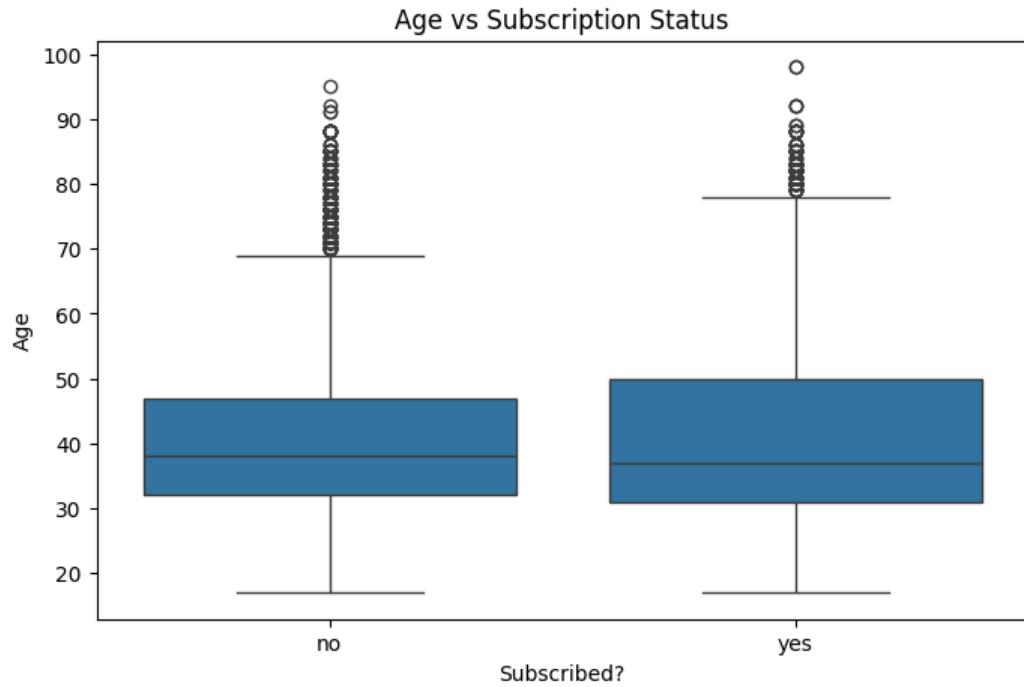
Age Distribution

```
In [29]: plt.figure(figsize=(8,4))
sns.histplot(df_cleaned['Age'], bins=20, kde=True)
plt.title("Age Distribution of Clients")
plt.xlabel("Age")
plt.ylabel("Count")
plt.show()
```



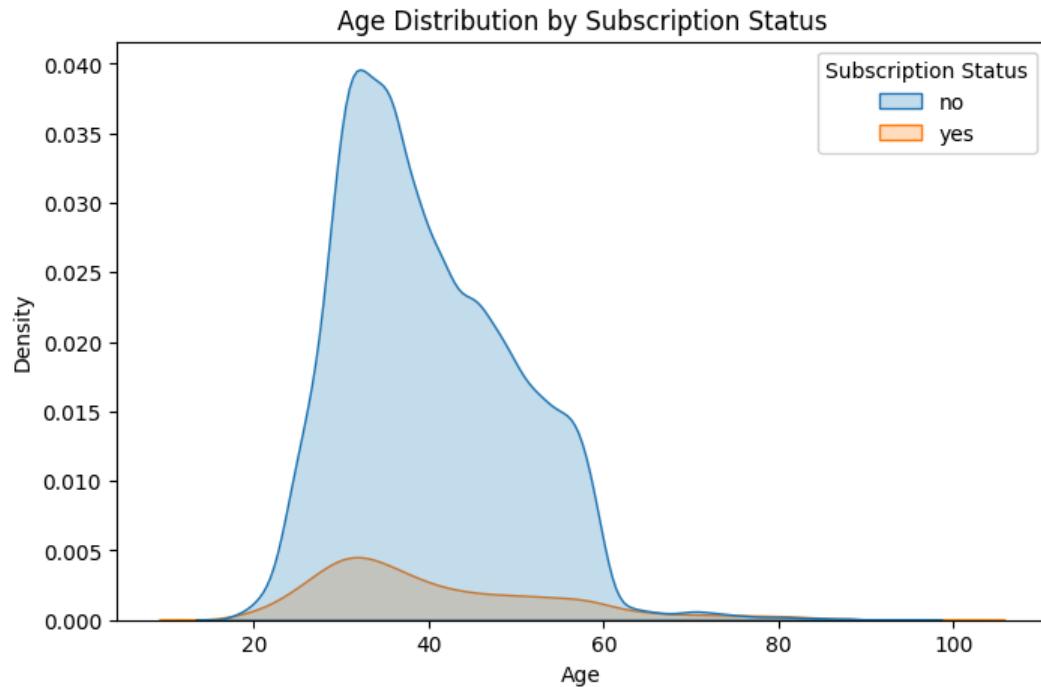
Comparing Age vs Subscription Status

```
In [30]: plt.figure(figsize=(8,5))
sns.boxplot(data=df_cleaned, x='Subscription Status', y='Age')
plt.title("Age vs Subscription Status")
plt.xlabel("Subscribed?")
plt.ylabel("Age")
plt.show()
```



The boxplot shows that the age distribution between customers who subscribed ("yes") and those who did not ("no") is very similar. Both groups have a median age around the late 30s. While there are older clients (shown as outliers), the overall spread of ages does not differ significantly between the two groups. This suggests that age alone is not a strong predictor of subscription.

```
In [31]: plt.figure(figsize=(8,5))
sns.kdeplot(data=df_cleaned, x='Age', hue='Subscription Status', fill=True)
plt.title("Age Distribution by Subscription Status")
plt.show()
```



The density plot shows that the age distributions for subscribers and non-subscribers are highly similar. Most clients fall between 25–60 years old, and both groups share the same general shape. This suggests that age is not a strong distinguishing factor in determining whether a client will subscribe. The dataset also contains very few clients above age 60.

Comparing Age Group vs Subscription Status

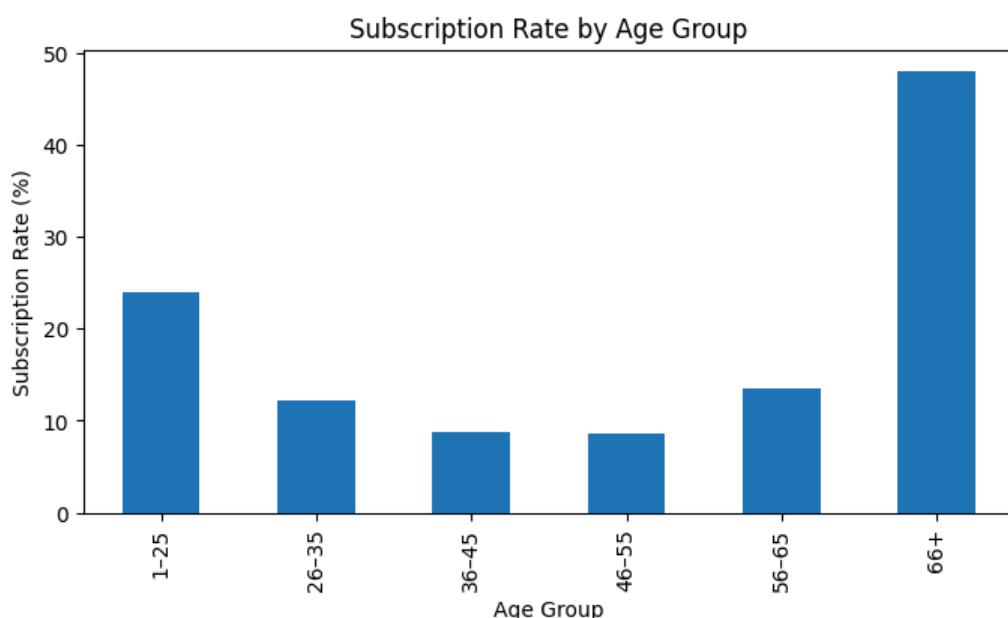
As individual ages data and analysis is very scarce, attempting to look at it in age groups would help to understand the data further

```
In [32]: # compute subscription rate (%) by age group
age_group_rates = (
    df_cleaned
        .groupby('Age Group')['Subscription Status']
        .value_counts(normalize=True)
        .unstack()
        .fillna(0)[['yes']] * 100
)
age_group_rates
```

C:\Users\jltho\AppData\Local\Temp\ipykernel_2920\3875872149.py:4: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
.groupby('Age Group')['Subscription Status']

```
Out[32]: Age Group
1-25      24.016563
26-35     12.253705
36-45      8.730420
46-55      8.600770
56-65     13.563004
66+       47.920133
Name: yes, dtype: float64
```

```
In [33]: plt.figure(figsize=(8,4))
age_group_rates.plot(kind='bar')
plt.title("Subscription Rate by Age Group")
plt.ylabel("Subscription Rate (%)")
plt.xlabel("Age Group")
plt.show()
```

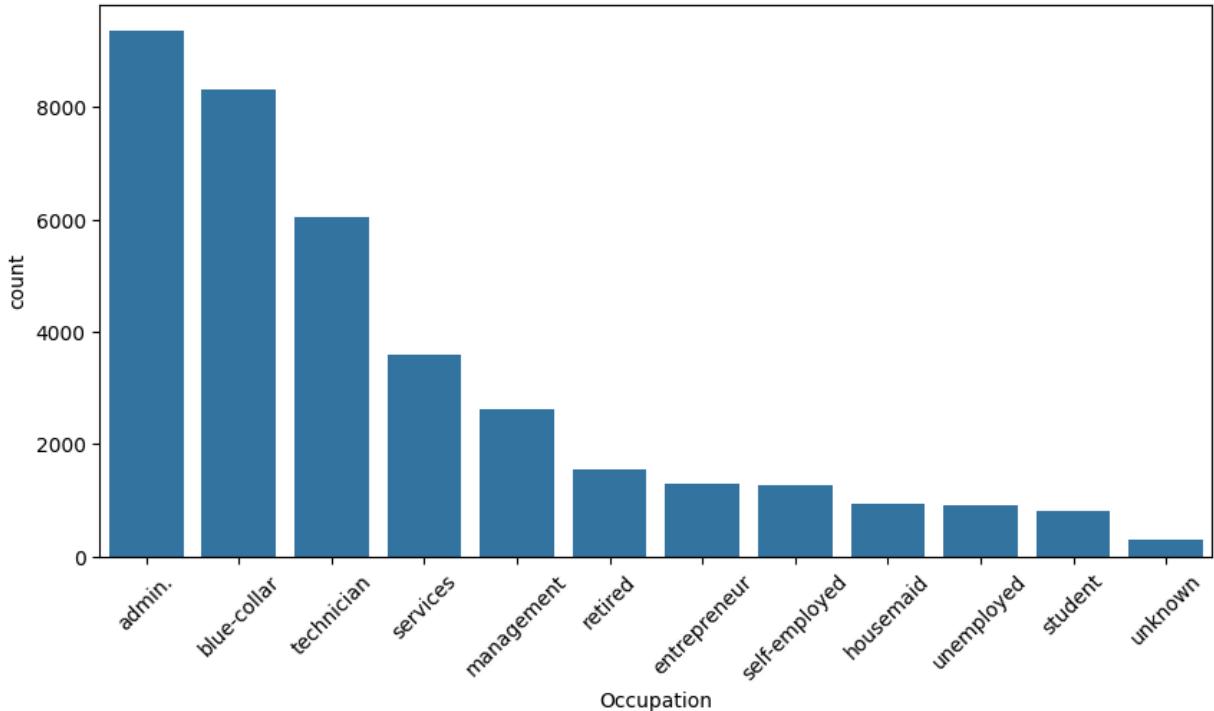


The clients most likely to subscribe are older adults (66+) and young adults (18–25), while middle-aged groups show significantly lower subscription rates, likely due to different financial priorities. This highlights clear age-based segmentation opportunities for future marketing strategies.

Occupations Distribution

```
In [34]: plt.figure(figsize=(10,5))
sns.countplot(data=df_cleaned, x='Occupation', order=df_cleaned['Occupation'].value_counts().index)
plt.title("Distribution of Occupations")
plt.xticks(rotation=45)
plt.show()
```

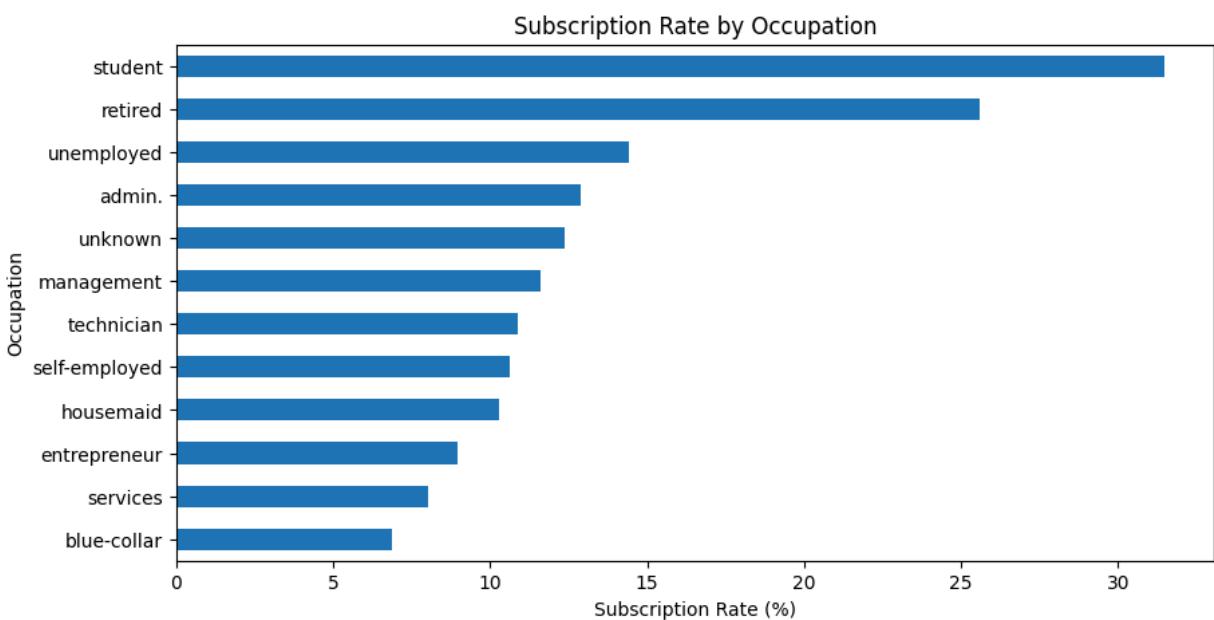
Distribution of Occupations



Subscription Rates vs Occupation

```
In [35]: occupation_rates = (
    df_cleaned.groupby('Occupation')['Subscription Status']
    .value_counts(normalize=True)
    .unstack()
    .fillna(0)['yes'] * 100
)

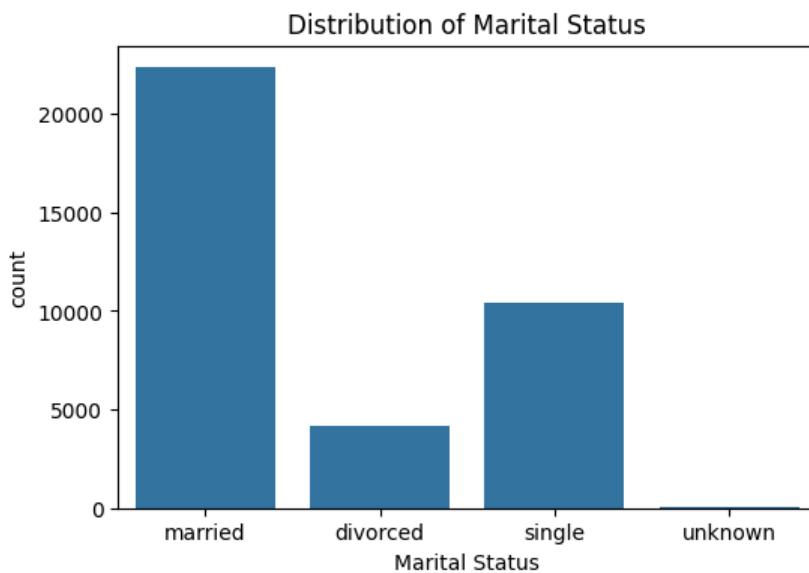
plt.figure(figsize=(10,5))
occupation_rates.sort_values().plot(kind='barh')
plt.title("Subscription Rate by Occupation")
plt.xlabel("Subscription Rate (%)")
plt.show()
```



Occupation has a clear impact on subscription behaviour. Students and retired clients show the highest subscription rates, making them strong target segments for term deposit campaigns. Middle-income groups such as admin, management, and technical roles show moderate response, while blue-collar and service-oriented occupations have the lowest subscription rates. This indicates that job stability, financial flexibility, and life stage significantly influence the likelihood of subscribing.

Distribution of Marital Status

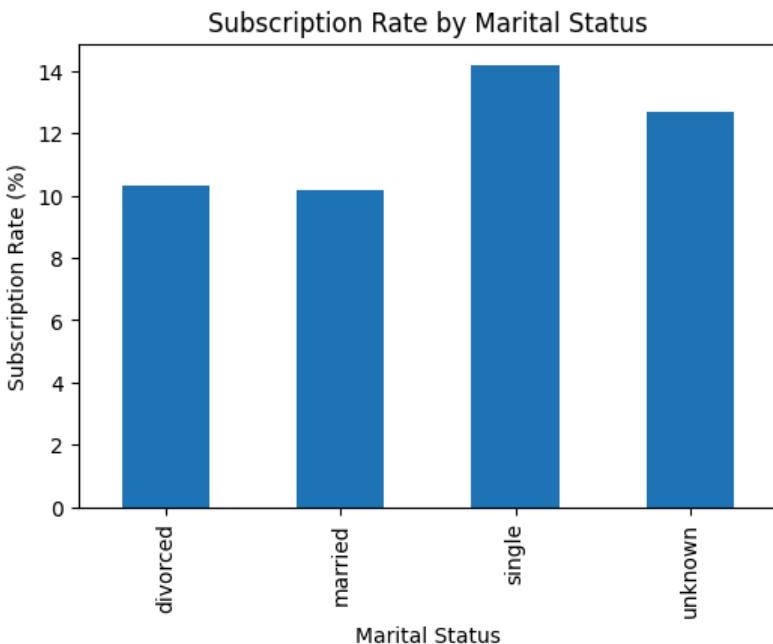
```
In [36]: plt.figure(figsize=(6,4))
sns.countplot(data=df_cleaned, x='Marital Status')
plt.title("Distribution of Marital Status")
plt.show()
```



Subscription Rate vs Marital Status

```
In [37]: marital_rates = (
    df_cleaned.groupby('Marital Status')['Subscription Status']
    .value_counts(normalize=True)
    .unstack()
    .fillna(0)['yes'] * 100
)

plt.figure(figsize=(6,4))
marital_rates.plot(kind='bar')
plt.title("Subscription Rate by Marital Status")
plt.ylabel("Subscription Rate (%)")
plt.show()
```

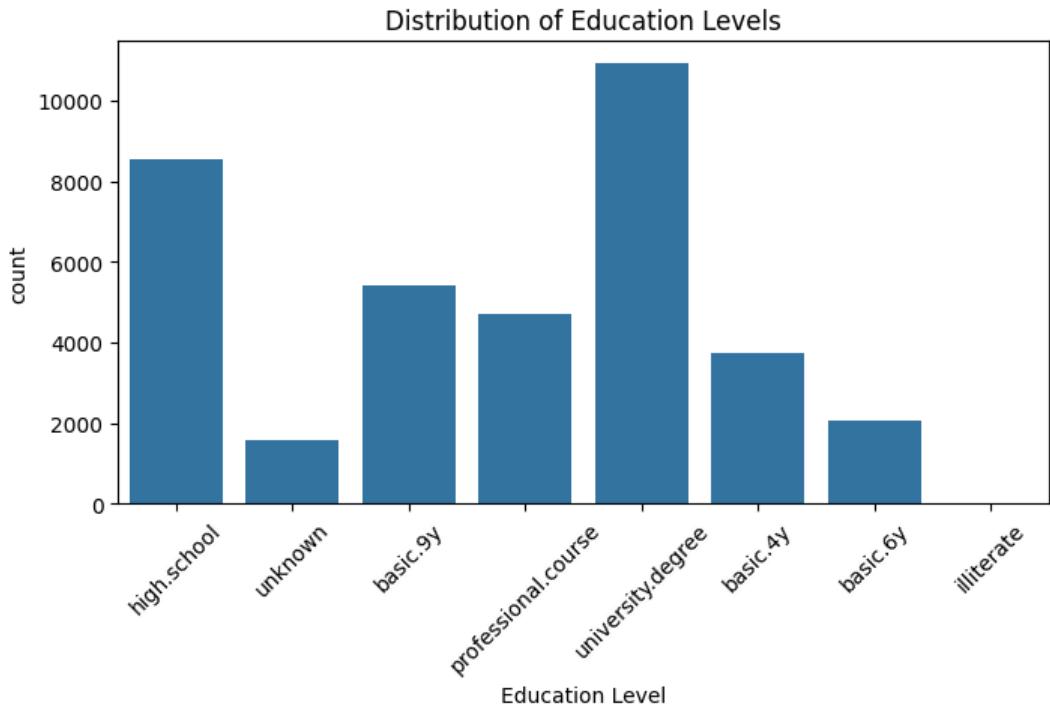


Marital status influences subscription behaviour. Single clients have the highest subscription rate, suggesting they are more receptive to term deposit offers. Married and divorced clients subscribe at similar moderate levels, likely due to greater financial commitments.

Interestingly, the "unknown" category also performs well, implying that missing marital data may correspond to younger or more financially active customers. Overall, marital status provides useful segmentation insights for targeted marketing.

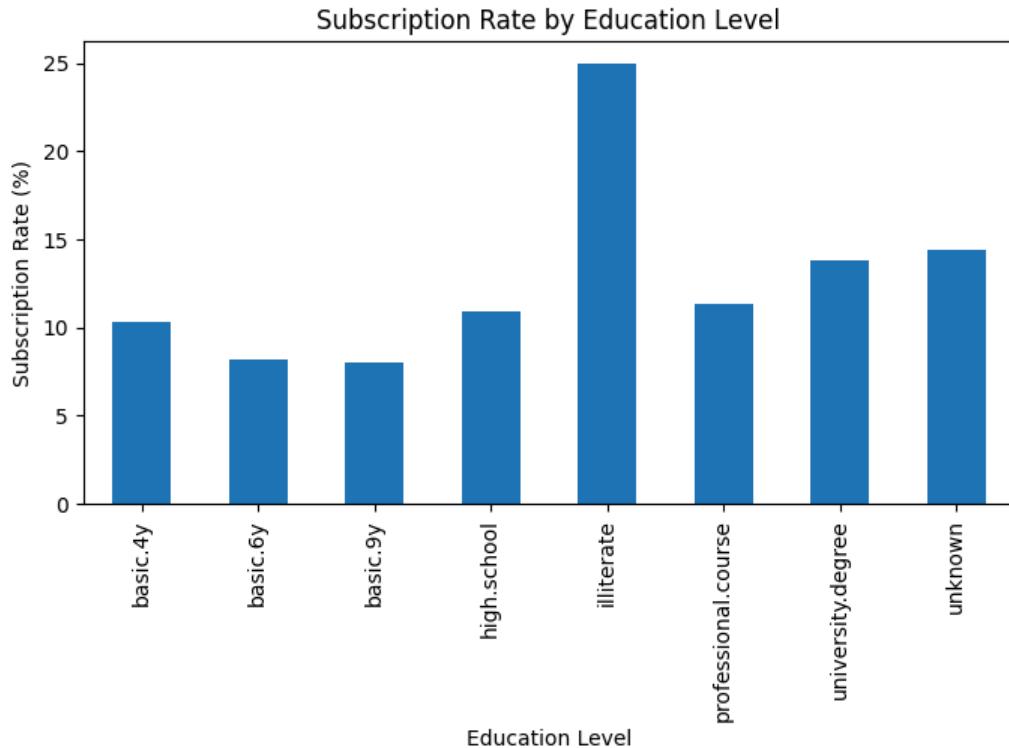
Education Levels Distribution

```
In [38]: plt.figure(figsize=(8,4))
sns.countplot(data=df_cleaned, x='Education Level')
plt.title("Distribution of Education Levels")
plt.xticks(rotation=45)
plt.show()
```



```
In [ ]: edu_rates = (
    df_cleaned.groupby('Education Level')['Subscription Status']
    .value_counts(normalize=True)
    .unstack()
    .fillna(0)['yes'] * 100
)

plt.figure(figsize=(8,4))
edu_rates.plot(kind='bar')
plt.title("Subscription Rate by Education Level")
plt.ylabel("Subscription Rate (%)")
plt.show()
```



Education level shows clear patterns in subscription behaviour. Clients with the highest education levels (university degree) tend to have higher subscription rates, likely due to greater financial literacy and income stability.

Surprisingly, the "illiterate" group has an extremely high subscription rate, but this is likely influenced by small sample size and should be interpreted with caution. Basic education groups consistently show the lowest subscription rates, indicating that financial understanding may influence willingness to invest in term deposits. Overall, education is a meaningful predictor of subscription behaviour.

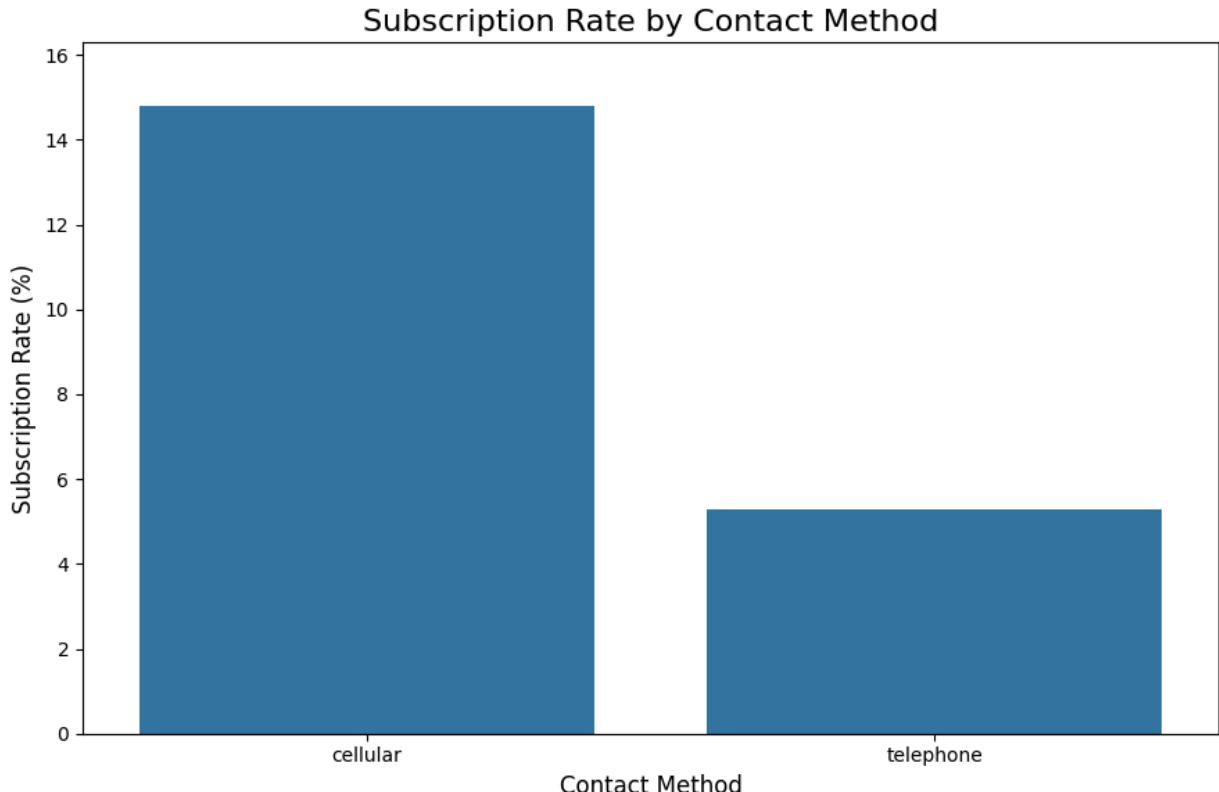
Contact with Clients

```
In [ ]: feature_col_cm = 'Contact Method'
print(f'{feature_col_cm} vs {target_col} Analysis')

#Calculation
contact_analysis = df_cleaned.groupby([feature_col_cm, target_col]).size().unstack(fill_value=0)
contact_analysis['Total'] = contact_analysis['no'] + contact_analysis['yes']
if 'yes' in contact_analysis.columns:
    contact_analysis['Subscription Rate (%)'] = (contact_analysis['yes'] / contact_analysis['Total']) * 100
else:
    contact_analysis['Subscription Rate (%)'] = 0
contact_analysis = contact_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(contact_analysis)

#Plotting
plt.figure(figsize=(9, 6))
sns.barplot(x=contact_analysis.index, y='Subscription Rate (%)', data=contact_analysis)
plt.title(f'Subscription Rate by {feature_col_cm}', fontsize=16)
plt.xlabel(f'{feature_col_cm}', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, contact_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('contact_method_subscription_rate.png')
plt.show()
```

Contact Method vs Subscription Status Analysis			
Subscription Status	no	yes	Total
Contact Method			
cellular	20017	3479	23496
telephone	12783	712	13495
			14.806776
			5.276028



The success rate for contacting clients via cellular phone is nearly three times higher (14.8%) than via telephone (5.23%).

Recommendation: Prioritize campaign resources, time, and budget towards clients for whom a cellular number is available.

```
In [ ]: # Campaign Calls
#Bins and Labels to store the thresholds to use
bins = [0, 1, 2, 3, 4, 5, 10, 60]
labels = ['1', '2', '3', '4', '5', '6-10', '11+']

df_cleaned['Campaign Bins'] = pd.cut(
    df_cleaned['Campaign Calls'],
    bins=bins,
    labels=labels,
    right=True,
    include_lowest=True
)

# Campaign Calls EDA
feature_col_cc = 'Campaign Bins'
print(f" Campaign Calls ({feature_col_cc}) vs {target_col} Analysis")

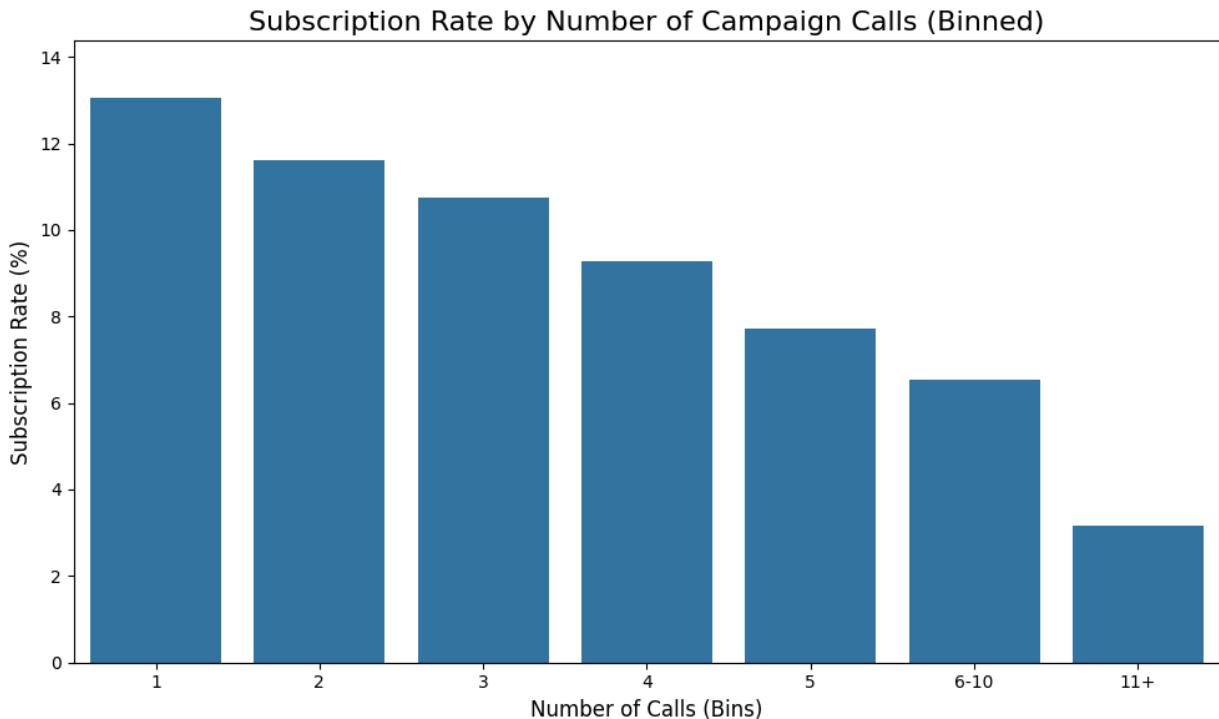
# Calculation
campaign_analysis = df_cleaned.groupby([feature_col_cc, target_col]).size().unstack(fill_value=0)
campaign_analysis['Total'] = campaign_analysis['no'] + campaign_analysis['yes']
if 'yes' in campaign_analysis.columns:
    campaign_analysis['Subscription Rate (%)'] = (campaign_analysis['yes'] / campaign_analysis['Total']) * 100
else:
    campaign_analysis['Subscription Rate (%)'] = 0
campaign_analysis = campaign_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(campaign_analysis)

# Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x=campaign_analysis.index, y='Subscription Rate (%)', data=campaign_analysis)
plt.title('Subscription Rate by Number of Campaign Calls (Binned)', fontsize=16)
plt.xlabel('Number of Calls (Bins)', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, campaign_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('campaign_calls_subscription_rate.png')
plt.show()
```

Campaign Status	no	yes	Total	Subscription Rate (%)
Campaign Bins				
1	13735	2064	15799	13.064118
2	8425	1107	9532	11.613512
3	4265	514	4779	10.755388
4	2181	223	2404	9.276206
5	1316	110	1426	7.713885
6-10	2114	148	2262	6.542882
11+	764	25	789	3.168568

```
C:\Users\jltho\AppData\Local\Temp\ipykernel_2920\907719962.py:19: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
```

```
campaign_analysis = df_cleaned.groupby([feature_col_cc, target_col]).size().unstack(fill_value=0)
```



The highest subscription rate comes from clients who were contacted only once (13%). The rate steadily decreases with each additional call. With 5 or more calls, the subscription rate drops significantly below 7%. Calls exceeding 10 are largely inefficient, resulting in a rate of just 3%.

```
In [ ]: #Previous Contact Days

# Indicate '999' means client was not contacted
df_cleaned['Was Previously Contacted'] = df_cleaned['Previous Contact Days'].apply(
    lambda x: 'Not Contacted' if x == 999 else 'Contacted'
)

# Previous Contact Days Category EDA
feature_col_pd = 'Was Previously Contacted'
print(f"Previous Contact Status ({feature_col_pd}) vs {target_col} Analysis")

# Calculation
pdays_category_analysis = df_cleaned.groupby([feature_col_pd, target_col]).size().unstack(fill_value=0)
pdays_category_analysis['Total'] = pdays_category_analysis['no'] + pdays_category_analysis['yes']
if 'yes' in pdays_category_analysis.columns:
    pdays_category_analysis['Subscription Rate (%)'] = (pdays_category_analysis['yes'] / pdays_category_analysis['Total'])
else:
    pdays_category_analysis['Subscription Rate (%)'] = 0
pdays_category_analysis = pdays_category_analysis.sort_values(by='Subscription Rate (%)', ascending=False)
print(pdays_category_analysis)

# Plotting
plt.figure(figsize=(7, 5))
sns.barplot(x=pdays_category_analysis.index, y='Subscription Rate (%)', data=pdays_category_analysis)
plt.title('Subscription Rate by Previous Contact Status (P-Days)', fontsize=16)
plt.xlabel('Previous Contact Status', fontsize=12)
plt.ylabel('Subscription Rate (%)', fontsize=12)
plt.ylim(0, pdays_category_analysis['Subscription Rate (%)'].max() * 1.1)
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig('pdays_category_subscription_rate.png')
```

```

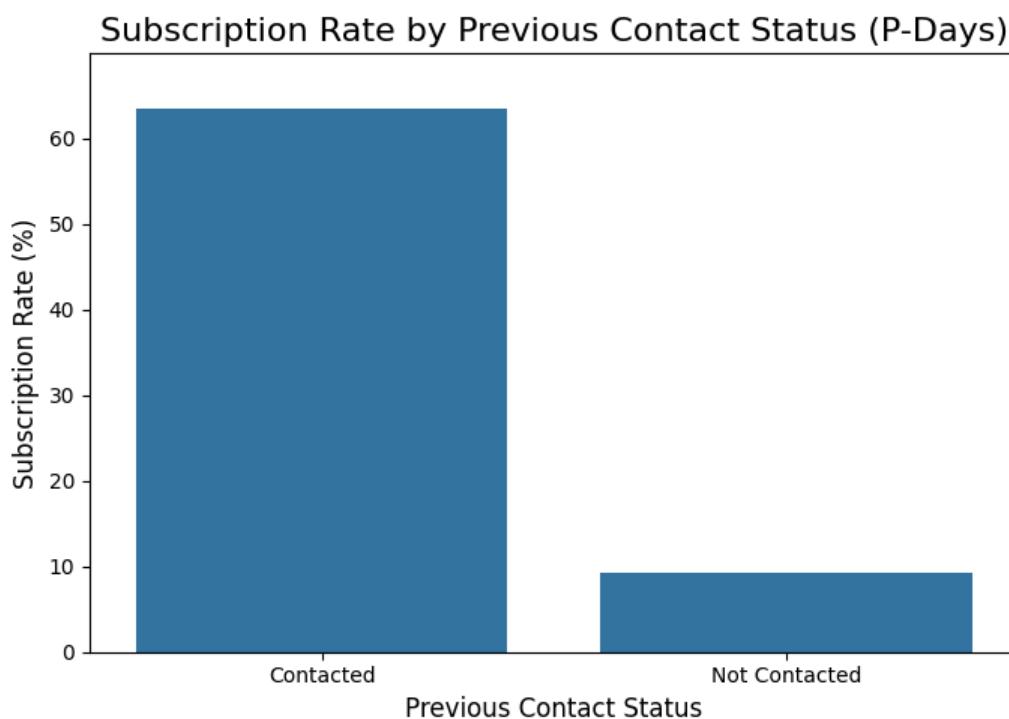
plt.show()

# Previous Contact Days for Contacted Clients (Mean)
print("\n Additional Insight: P-Days for Contacted Clients (Mean)")
contacted_df = df_cleaned[df_cleaned['Was Previously Contacted'] == 'Contacted']
pdays_insight = contacted_df.groupby('Subscription Status')['Previous Contact Days'].mean()
print(pdays_insight)

```

Previous Contact Status (Was Previously Contacted) vs Subscription Status Analysis

Subscription Status	no	yes	Total	Subscription Rate (%)
Was Previously Contacted				
Contacted	501	872	1373	63.510561
Not Contacted	32299	3319	35618	9.318322



```

Additional Insight: P-Days for Contacted Clients (Mean)
Subscription Status
no      6.083832
yes     5.891055
Name: Previous Contact Days, dtype: float64

```

Clients who were successfully contacted in a previous campaign (clients with 999 value) have an overwhelmingly high subscription rate of 63.83%. This indicates that clients who engaged previously are highly receptive.

Average P-Days for Contacted Clients: Subscribed ('yes'): 5.9 days since last contact. Not Subscribed ('no'): 6.1 days since last contact. The difference is negligible, suggesting previous contact with the client matters more than how long ago that contact was (at least within the 6-day average window).